

Trabalho 04

Do livro AMP, devem ser feitos (no mínimo!) os seguintes exercícios:
Capítulo 2. 11, 13, 15, 16

Exercise 11. Programmers at the Flaky Computer Corporation designed the protocol shown in Fig. 2.15 to achieve n-thread mutual exclusion. For each question, either sketch a proof, or display an execution where it fails.

```
1 class Flaky implements Lock {
2     private int turn;
3     private boolean busy = false;
4     public void lock() {
5         int me = ThreadID.get();
6         do {
7             do {
8                 turn = me;
9             } while (busy);
10            busy = true;
11        } while (turn != me);
12    }
13    public void unlock() {
14        busy = false;
15    }
16 }
```

Figure 2.15 The Flaky lock used in Exercise 11.

- Does this protocol satisfy mutual exclusion?

Seguindo a notação de *prova* aplicada no livro, dada uma execução de duas Threads A e B, onde A executa primeiro que B, é possível dizer que o mesmo satisfaz a exclusão mútua (mutex)

$\text{write}_A(\text{turn} = \text{me}(A)) \rightarrow \text{read}_A(\text{busy} == \text{false}) \rightarrow \text{write}_A(\text{busy} == \text{true}) \rightarrow \text{read}_A(\text{turn} == \text{me}(A)) \rightarrow \text{CS}_A$
 $\text{write}_B(\text{turn} = \text{me}(B)) \rightarrow \text{read}_B(\text{busy} == \text{false}) \rightarrow \text{write}_B(\text{busy} == \text{true}) \rightarrow \text{read}_B(\text{turn} == \text{me}(B)) \rightarrow \text{CS}_B$

$\text{read}_A(\text{turn} == \text{me}(A)) \rightarrow \text{write}_B(\text{turn} = \text{me}(B))$ //Exclusão Mútua

Como é usado *Do... While* é testado primeiro quem é *turn*. A Thread A e B não acessaram a área crítica ao mesmo tempo.

- Is this protocol starvation-free?

Não, pelo fato de protocolos que podem ser deadlock não serem starvation-free.

Se uma thread está tentando entrar em sua seção crítica, então a thread deve eventualmente entrar em sua seção crítica, no exemplo a Thread A pode não entrar em sua região crítica.

```

#Thread A: writeA(turn = me(A))
#Thread B: writeB(turn = me(B))
#Thread B: readB(busy == false) → writeB(busy == true) → readB(turn == me (B)) → CSB
#Thread A: writeA(turn = me(A)) → readA(busy == false)
#Thread B: CSB → writeB(busy == true)
#Thread B: writeB(turn = me(B))

```

```

1  class Flaky implements Lock {
2      private int turn;
3      private boolean busy = false;
4      public void lock() {
5          int me = ThreadID.get();
6          do {
7              do {
8                  turn = me;
9              } while (busy);
10             busy = true;
11         } while (turn != me);
12     }
13     public void unlock() {
14         busy = false;
15     }
16 }

```

Figure 2.15 The Flaky lock used in Exercise 11.

- Is this protocol deadlock-free?

O loop interno continua a ser executado, já que $busy == true$. Como mencionado no Capítulo 1 a exclusão mútua exige deadlock-free (página 11).

```

#Thread A: writeA(turn = me(A))
#Thread B: writeB(turn = me(B))
#Thread B: readB(busy == false) → writeB(busy == true)
#Thread A: writeA(turn = me(A))
#Thread B: readB(busy == false) → writeB(busy == true)

```

Exercise 13. Another way to generalize the two-thread Peterson lock is to arrange a number of 2-thread Peterson locks in a binary tree. Suppose n is a power of two. Each thread is assigned a leaf lock which it shares with one other thread. Each lock treats one thread as thread 0 and the other as thread 1.

In the tree-lock's acquire method, the thread acquires every two-thread Peterson lock from that thread's leaf to the root. The tree-lock's release method for the tree-lock unlocks each of the 2-thread Peterson locks that thread has acquired, from the root back to its leaf. At any time, a thread can be delayed for a finite duration. (In other words, threads can take naps, or even vacations, but they do not drop dead.)

For each property, either sketch a proof that it holds, or describe a (possibly infinite) execution where it is violated.

1. mutual exclusion.
2. freedom from deadlock.
3. freedom from starvation.

Is there an upper bound on the number of times the tree-lock can be acquired and released between the time a thread starts acquiring the tree-lock and when it succeeds?

Exercise 15. In practice, almost all lock acquisitions are uncontended, so the most practical measure of a lock's performance is the number of steps needed for a thread to acquire a lock when no other thread is concurrently trying to acquire the lock.

Scientists at Cantaloupe-Melon University have devised the following “wrapper” for an arbitrary lock, shown in Fig. 2.16. They claim that if the base Lock class provides mutual exclusion and is starvation-free, so does the FastPath lock, but it can be acquired in a constant number of steps in the absence of contention. Sketch an argument why they are right, or give a counterexample.

```
1 class FastPath implements Lock {
2   private static ThreadLocal<Integer> myIndex;
3   private Lock lock;
4   private int x, y = -1;
5   public void lock() {
6     int i = myIndex.get();
7     x = i;                               // I'm here
8     while (y != -1) {}                  // is the lock free?
9     y = i;                               // me again?
10    if (x != i)                          // Am I still here?
11      lock.lock();                       // slow path
12  }
13  public void unlock() {
14    y = -1;
15    lock.unlock();
16  }
17 }
```

Figure 2.16 Fast path mutual exclusion algorithm used in Exercise 15.

Line 4: $\text{write}_A(x = A) \rightarrow \text{write}_B(x = B) \rightarrow$

Line 8: $\text{read}_A(y == A) \rightarrow \text{read}_B(y == B) \rightarrow$

Line 4: $\text{write}_C(x = C)$

Line 10: $\text{read}_A(x == A) \rightarrow \text{read}_B(x == A)$ // Os dois podem entrar no lock ao mesmo tempo

Como não há garantia de exclusão mútua duas Threads podem ficar impedidas de continuar a execução e o acesso a CS pode nunca acontecer. Então não é starvation-free onde se um thread está tentando entrar em sua seção crítica, então esta thread deve eventualmente entrar em sua seção crítica.

Exercise 16. Suppose n threads call the `visit()` method of the Bouncer class shown in Fig. 2.17. Prove that:

- At most one thread gets the value STOP.
- At most $n - 1$ threads get the value DOWN.
- At most $n - 1$ threads get the value RIGHT.

Note that the last two proofs are not symmetric.

```
1  class Bouncer {
2      public static final int DOWN = 0;
3      public static final int RIGHT = 1;
4      public static final int STOP = 2;
5      private boolean goRight = false;
6      private ThreadLocal<Integer> myIndex;
7      private int last = -1;
8      int visit() {
9          int i = myIndex.get();
10         last = i;
11         if (goRight)
12             return RIGHT;
13         goRight = true;
14         if (last == i)
15             return STOP;
16         else
17             return DOWN;
18     }
19 }
```

Figure 2.17 The Bouncer class implementation.

[...]